

PHANTOM — Updated Results Figures

Per-window M1 · 60% penetration · seed 1 (A4 = 3-seed mean)

Generated 2026-10-02 · all figures A1/A2 fusion-fix applied

- Fig 1 SOTA comparison (A1-A4)
- Fig 2 Fairness reversal (A2/A4)
- Fig 3 A2 fusion fix
- Fig 4 Exp1 penetration x intensity
- Fig 5 Exp4 attack selectivity
- Fig 6 Exp2 mobility / speed
- Fig 7 Exp3 network scale (N=100/200/300)

Exp3 MCC curve crypto-off (MCC-neutral); N=400 excluded (crash pending).

Fig 1 — PHANTOM vs fairly-scored SOTA baselines




Fig 2 — Baselines' wins were a scoring artifact (latched → per-window)




Fig 3 — A2 fusion fix recovers its S1+S2 primary




Fig 4 — Exp1: detection vs penetration × intensity




Fig 5 — Exp4: detection vs attack selectivity (AOEI)




Fig 6 — Exp2: detection vs mobility (fix-applied)




Fig 7 — Exp3: detection vs scale, N=100/200/300 (fix-applied)

