

PHANTOM — Updated Results Figures

Per-window M1 scoring · 60% penetration · seed 1 (A4 = 3-seed mean)


Generated 2026-10-02 · commits 1c12582 / 0c3b2d7 / 60f7a2a

Contents:

- Fig 1 SOTA comparison (A1-A4)
- Fig 2 Fairness reversal (A2/A4)
- Fig 3 A2 fusion fix
- Fig 4 Exp1 trend (penetration x intensity)
- Fig 5 Exp4 trend (selectivity)


Excluded (not current): Exp2 speed and Exp3 scale — pre-fix (09-19);
Exp3 N=400 cell still pending a crash fix.

Fig 1 - SOTA comparison (per-window MCC): PHANTOM vs fairly-scored baselines




A2 0.920 vs TAP 0.60; A4 0.714 (3-seed) vs SFTO 0.55; A3 tie; A1 0.647 (FPR = attribution artifact).

Fig 2 - Fairness reversal: baselines' latched claim collapses under per-window M1




TAP A2 0.81->0.60, SFTO A4 0.96->0.55. The apparent advantage was the scoring convention, not the detector.

Fig 3 - A2 fusion fix (our own bug): composite recovers its S1+S2 primary




Covering-RSU attribution + S1 fold-in: A2 composite 0.594 -> 0.822 -> 0.920, FPR 33.6% -> 7%.

Fig 4 - Exp1: detection MCC vs penetration x intensity (A1-A4), fix-consistent



A2 climbs 0.77 (p20) -> 0.89 (p40) -> 0.91 (p60), consistent with the SOTA table.

Fig 5 - Exp4: detection MCC vs attack selectivity (AOEI)



Lower target ratio = stealthier attack; PHANTOM degrades gracefully.